


Conjugate gradients: ingredients


- Quadratic form
- Steepest descent
- Eigenvectors and Eigenvalues
- Conjugate directions

Quadratic forms of a matrix


Positive definite




Negative definite



Singular, positive definite



Indefinite






Quadratic forms

General quadratic equation

$$f(x) = \frac{1}{2} x^T A x - x^T b$$

$f(x)$ is minimized by the solution to $Ax=b$, which is demonstrated on the next slide.



Gradient points in the direction of steepest increase of $f(x)$.

Conjugate gradient (cont'd)

Gradient

$$f'(x) = \begin{bmatrix} \frac{\partial}{\partial x_1} f(x) \\ \frac{\partial}{\partial x_2} f(x) \\ \vdots \\ \frac{\partial}{\partial x_n} f(x) \end{bmatrix}$$

With some more math we arrive at

$$f'(x) = \frac{1}{2} x^T A + \frac{1}{2} (Ax)^T - b^T$$

which, if A is symmetric, reduces to

$$f'(x) = Ax - b$$

Setting $f'(x)=0$, this is equation we need to solve (at the minimum, the first derivative is zero).

Thus, $Ax=b$ can be solved by finding x that minimizes $f(x)$. This holds, if A is symmetric and positive definite!

Steepest descent

Question: How to get to the minimum (fast, efficiently)?

This is a general question, which arises in many different contexts, such as solving PDEs and inversion.

Suggestion: slide down to the minimum from an arbitrary starting point x_0 following a number of steps, x_1, x_2, \dots, x_n until we are close enough to the solution x .

$-f'(x_i) = b - Ax_i$ This would be the direction to follow in which $f(x)$ decreases fastest (direction of steepest descent).

$e_i = x_i - x$ This is the **error vector**, which indicates the distance from the true solution vector, x , at each step, $i=1, \dots, n$.

$r_i = b - Ax_i$ This is the **residual vector**, which indicates the distance from the true b vector at each step, $i=1, \dots, n$. Think of it as
 $r_i = -f'(x_i)$ the error vector transformed by A into the b -space. The residual vector is actually the direction of the steepest descent!

Steepest descent (cont'd)

These are then the different steps of the Steepest Descent method:

$$r_i = b - Ax_i$$


$$\alpha_i = \frac{r_i^T r_i}{r_i^T A r_i}$$

$$x_{i+1} = x_i + \alpha_i r_i$$

The method requires two matrix-vector products per iteration! One can be eliminated though:

$$r_{i+1} = r_i - \alpha_i A r_i$$

Remember i and $i+1$ are the iteration counts of the steepest descent solver (NOT spatial indices)




Why is there a distinct zigzag path? Remember that α should be chosen that r_0 and $f'(x_1)$ are orthogonal.

Convergence of Steepest Descent: Instant convergence (cont'd)

There is only one step to the exact solution, because the x_i lies exactly on one axis of the ellipsoid and the residual points directly into the center of the ellipsoid;
thus, $\alpha_i = 1/\lambda_e$ results in instant convergence.

This is a very special case.

For a more general analysis we again think of a vector as a sum of other well-understood vectors i.e. express e_i as a linear combination of *orthonormal* eigenvectors.



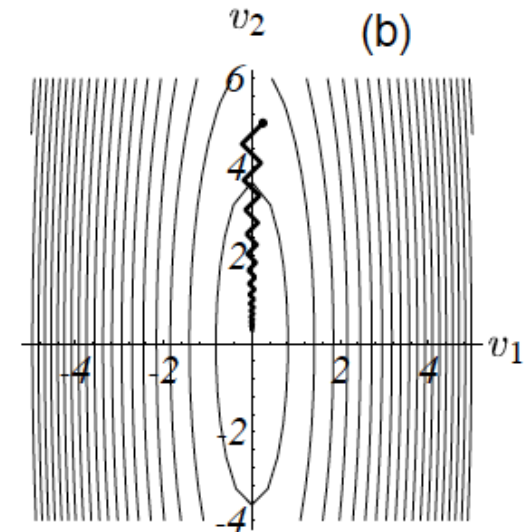

Convergence of Steepest Descent: General convergence (cont'd)

In the suma, for demonstration we have $n = 2$ and assume $\lambda_1 \geq \lambda_2$.


$$\kappa = \frac{\lambda_1}{\lambda_2} \geq 1 \quad \text{spectral condition number}$$

$$\mu = \frac{\xi_2}{\xi_1} \quad \text{slope of } e_i$$

$$\begin{aligned} \omega^2 &= 1 - \frac{(\xi_1^2 \lambda_1^2 + \xi_2^2 \lambda_2^2)^2}{(\xi_1^2 \lambda_1 + \xi_2^2 \lambda_2)(\xi_1^2 \lambda_1^3 + \xi_2^2 \lambda_2^3)} \\ &= 1 - \frac{(\kappa^2 + \mu^2)^2}{(\kappa + \mu^2)(\kappa^3 + \mu^2)} \end{aligned}$$




Convergence of Steepest Descent: General convergence (cont'd)




The more ill-conditioned the matrix is i.e. the larger the condition number the slower is the convergence.

What is a disadvantage of Steepest Descent?



Well, there are multiple steps in the same direction of various length.

What if we could find a set of orthogonal search directions d_0, d_1, \dots, d_n where we take exactly one step in each direction!





Conjugate Directions (cont'd)

Solution: make the search directions A -orthogonal or conjugate instead of orthogonal

$$v_i^T A v_j = 0 \quad (\text{here } v_i \text{ could be the direction, } d, \text{ and } v_j \text{ could be some error, } e).$$

These vectors are A -orthogonal, because...they are orthogonal after being “stretched”.



Gram-Schmidt Conjugation

The last thing needed is a set of A -orthogonal search directions.


Suppose a set of n linearly independent vectors $u_0, u_1, u_2, \dots, u_n$ (could be the coordinate axes). To construct d_i , take u_i and subtract any components that are not A -orthogonal. Thus, set $d_0 = u_0$ and for $i > 0$

$$d_0 = u_0$$

$$d_i = u_i - \sum_{j=0}^{i-1} \beta_{ij} d_j$$

$$\beta_{ij} = \frac{u_i^T A d_j}{d_j^T A d_j}$$

Note, there is an outer loop over $j \leq n$ search directions spanning the whole vector space!



Begin with two linearly independent vectors u_0, u_1 . Set $d_0 = u_0$. The vector u_1 is composed of two components: u^* (A -orthogonal to d_0) and u^+ (parallel to d_0). After conjugation only A -orthogonal portion remains and $d_1 = u^*$.

The disadvantage is, all old search vectors must be kept (in memory) to construct the new one and $O(n^3)$ operations are required!

Conjugate gradients (cont'd)

Putting it all together:

$$d_0 = r_0 = b - Ax_0$$


$$\alpha_i = \frac{r_i^T r_i}{d_i^T A d_i}$$

$$x_{i+1} = x_i + \alpha_i d_i$$

$$r_{i+1} = r_i - \alpha_i A d_i$$

$$\beta_{i+1} = \frac{r_{i+1}^T r_{i+1}}{r_i^T r_i}$$

$$d_{i+1} = r_{i+1} + \beta_{i+1} d_i$$



Conjugate gradients (cont'd)

Putting it all together:

$$d_0 = r_0 = b - Ax_0$$

$$\alpha_i = \frac{r_i^T r_i}{d_i^T A d_i}$$

$$x_{i+1} = x_i + \alpha_i d_i$$

$$r_{i+1} = r_i - \alpha_i A d_i$$

$$\beta_{i+1} = \frac{r_{i+1}^T r_{i+1}}{r_i^T r_i}$$

$$d_{i+1} = r_{i+1} + \beta_{i+1} d_i$$

Given the inputs A , b , x_0 , maximum iteration number i_{max} , and error tolerance $\varepsilon < 1$:

```
i ← 0
r ← b - Ax
d ← r
δnew ← rTr
δ0 ← δnew
While i < imax and δnew > ε2δ0 do
    q ← Ad
    α ← δnew / (dTq)
    x ← x + αd
    If i is divisible by 50
        r ← b - Ax
    else
        r ← r - αq
    δold ← δnew
    δnew ← rTr
    β ← δnew / δold
    d ← r + βd
    i ← i + 1
```

Preconditioning (cont'd)

Untransformed CG Method

$$r_0 = b - Ax_0$$

$$d_0 = M^{-1}r_0$$

$$\alpha_i = \frac{r_i^T M^{-1} r_i}{d_i^T A d_i}$$

$$x_{i+1} = x_i + \alpha_i d_i$$

$$r_{i+1} = r_i - \alpha_i A d_i$$



$$\beta_{i+1} = \frac{r_{i+1}^T M^{-1} r_{i+1}}{r_i^T M^{-1} r_i}$$

$$d_{i+1} = M^{-1} r_{i+1} + \beta_{i+1} d_i$$

Preconditioning (cont'd)

What does preconditioning due to the quadratic form? It stretched it and makes it appear more spherical.

There is a (very) large number of preconditioners out there and many consider it some kind of “black art”. One example is simply a diagonal matrix with entries taken from A , which scales the quadratic form along the coordinate axes. Note, in case of the large problems, application of CG generally always requires a preconditioner of some sorts.



Concrete example of how to solve a set of non-linear system of equations

Additional assumptions:

1. We neglect gravity (horizontal flow only) in 1D
2. $C(p) = \text{const.}$ and $D = K / C = \text{const.}$

$$\frac{\partial p_w}{\partial t} = D \frac{\partial}{\partial x} \left[k_w(p_w) \frac{\partial p_w}{\partial x} \right]$$

Finite difference, implicit

$$\frac{p_i^{k+1} - p_i^k}{\Delta t} = \frac{D}{\Delta x^2} \left[k_{uw}(p) (p_{i-1}^{k+1} - p_i^{k+1}) + k_{uw}(p) (p_{i+1}^{k+1} - p_i^{k+1}) \right]$$

Instead of using the harmonic mean for the relative permeabilities $k(p)$, we upwind $k_{uw}(p)$ i.e. take the values from the cell where the flow is coming from!

Concrete example and how to solve a set of non-linear equation

We can re-write the equation

$$0 \equiv F(p^k) = p_i^{k+1} - p_i^k - \frac{D\Delta t}{\Delta x^2} \left[k_{uw}(p) (p_{i-1}^{k+1} - p_i^{k+1}) + k_{uw}(p) (p_{i+1}^{k+1} - p_i^{k+1}) \right]$$

$$0 \equiv F(p^{k+1}) = p^{k+1} - p^k - \mathbf{C}f(p^{k+1})$$

\mathbf{C} is a matrix containing the constant coefficients, $D\Delta t / \Delta x^2$.

\mathbf{C} is not the \mathbf{A} matrix from our $\mathbf{Ax} = \mathbf{b}$ problem!

While \mathbf{C} is a sparse matrix, its shape differs from \mathbf{A} .

Applying the Newton-Raphson method

$$0 \equiv F(p^{k+1}) = p^{k+1} - p^k - \mathbf{C}f(p^{k+1})$$

Because of the nonlinearity, we only achieve an approximation \tilde{p}^{k+1} of the true p^{k+1} with one linear solve of the above system of equations. But we may be able to improve the approximation in an iterative fashion.

Using (again) a Taylor expansion we have

$$F(p^{k+1}) \approx F(\tilde{p}^{k+1}) + F'(\tilde{p}^{k+1})(p^{k+1} - \tilde{p}^{k+1})$$

An improved approximation of p^{k+1} is obtained by

$$\tilde{p}^{k+1} + d\tilde{p}^{k+1} - F'(\tilde{p}^{k+1})d\tilde{p}^{k+1} = F(\tilde{p}^{k+1})$$

Solve the system of equations to find the update

$$\underbrace{d\tilde{p}^{k+1}}_{\text{update}} = -\underbrace{F'(\tilde{p}^{k+1})^{-1}}_{\text{Jacobi matrix}} \underbrace{F(\tilde{p}^{k+1})}_{\text{Approx. solution vector}}$$

Applying the Newton-Raphson method (cont'd)

The Jacobian

$$F'(p^{k+1}) = J(p^{k+1}) = \frac{\partial(p^{k+1} - p^k)}{\partial p^{k+1}} - \mathbf{C} \frac{\partial f(p^{k+1})}{\partial p^{k+1}} = 1 - \mathbf{C} \frac{\partial f(p^{k+1})}{\partial p^{k+1}}$$

$$J = \begin{bmatrix} \frac{\partial F_1}{\partial p_1} & \frac{\partial F_1}{\partial p_2} & \dots & \\ \frac{\partial F_2}{\partial p_1} & \frac{\partial F_2}{\partial p_2} & \frac{\partial F_2}{\partial p_3} & \\ \vdots & \vdots & \ddots & \vdots \\ \dots & \dots & \dots & \frac{\partial F_N}{\partial p_N} \end{bmatrix}$$

What happens if $k(p) = k = \text{const.}$ in $f(p^{k+1})$?

And how does J look in case of 2D and 3D?

Applying the Newton-Raphson method (cont'd)

$$F_1 = p_1^{k+1} - p_1^k - \frac{D\Delta t}{\Delta x^2} \left[k_{uw}(p_1)(p_0^{k+1} - p_1^k) + k_{uw}(p_2)(p_2^{k+1} - p_1^{k+1}) \right]$$

$$\frac{\partial F_1}{\partial p_2} = -\frac{D\Delta t}{\Delta x^2} \left[\frac{\partial k_{uw}(p_2)}{\partial p_2} (p_2 - p_1) + k_{uw}(p_2) \right]$$

$$F_2 = p_2^{k+1} - p_2^k - \frac{D\Delta t}{\Delta x^2} \left[k_{uw}(p_2)(p_1^{k+1} - p_2^{k+1}) + k_{uw}(p_3)(p_3^{k+1} - p_2^{k+1}) \right]$$

$$\frac{\partial F_2}{\partial p_1} = -\frac{D\Delta t}{\Delta x^2} [-k_{uw}(p_2)]$$

Etc.

Each row contains the transposed gradients of the function f .

Note, the Jacobian can also be approximated via a Taylor expansion (differencing), which however requires a large number of non-linear function evaluations.

Applying the Newton-Raphson method (cont'd)

The different steps of the Newton Raphson are

1. Calculate $F(p^n)$ starting e.g. from an initial guess of p , where n is the number of Newton-Raphson iteration.
2. Calculate the derivative $F'(p^n)$ i.e. the Jacobi matrix
3. The improvement dp^n is calculated as $dp^n = -F'(p^n)^{-1}F$
4. Calculate $p^{n+1} = p^n + dp^n$
5. Repeat steps 1 to 4 until $norm(dp^n)$ has decreases below a pre-defined limit

MATLAB/pseudo-code snipped:

```
niter = 0
while (niter < nmax & err > toll),
    niter = niter+1;
    [relperm,drelperm] = my_relperm(...)
    fp = my_fp()
    F = p - p_old - C*fp; % approx. soln vec
    dF = my_jacobian(...); % Analytical J or FD
    dp = -dF\F; %Solve for the update dp
    err = abs(dp); % Is this correct?
    p = p+dp
end
```

Thus, the Newton-Raphson method combines an outer non-linear iteration method with an inner linear systems solver approach.

The basic approach makes up the large class of Newton-PCG solvers, where a Newton method is combined with a PCG linear solver.